

Amazon Product Review: Sentiment Analysis and Star Rating Prediction

Install required libraries

```
!pip install swifter
!pip install pyLDAvis
!pip install unicode
!pip install TextBlob
!pip install text2emotion
```



Collecting swifter

Downloading <https://files.pythonhosted.org/packages/f4/3b/04bf42b94a22725241b47e025>
|████████████████████| 634kB 7.5MB/s

Requirement already satisfied: pandas>=1.0.0 in /usr/local/lib/python3.6/dist-package
Collecting psutil>=5.6.6

Downloading <https://files.pythonhosted.org/packages/33/e0/82d459af36bda999f82c7ea86>
|████████████████████| 471kB 14.7MB/s

Requirement already satisfied: dask[dataframe]>=2.10.0 in /usr/local/lib/python3.6/di
Requirement already satisfied: tqdm>=4.33.0 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: ipywidgets>=7.0.0cloudpickle>=0.2.2 in /usr/local/lib/
Requirement already satisfied: parso>0.4.0 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: bleach>=3.1.1 in /usr/local/lib/python3.6/dist-package
Collecting modin[ray]>=0.8.1.1

Downloading <https://files.pythonhosted.org/packages/ab/a9/ead212fa94de8f14459e22b06>
|████████████████████| 542kB 20.1MB/s

Requirement already satisfied: python-dateutil>=2.7.3 in /usr/local/lib/python3.6/dis-
Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: numpy>=1.15.4 in /usr/local/lib/python3.6/dist-package
Collecting partd>=0.3.10; extra == "dataframe"

Downloading <https://files.pythonhosted.org/packages/44/e1/68dbe731c9c067655bff1eca5>

Requirement already satisfied: toolz>=0.7.3; extra == "dataframe" in /usr/local/lib/p
Collecting fsspec>=0.6.0; extra == "dataframe"

Downloading <https://files.pythonhosted.org/packages/a5/8b/1df260f860f17cb0869817015>
|████████████████████| 92kB 9.4MB/s

Requirement already satisfied: ipython>=4.0.0; python_version >= "3.3" in /usr/local/
Requirement already satisfied: nbformat>=4.2.0 in /usr/local/lib/python3.6/dist-packa
Requirement already satisfied: traitlets>=4.3.1 in /usr/local/lib/python3.6/dist-pack
Requirement already satisfied: ipykernel>=4.5.1 in /usr/local/lib/python3.6/dist-pack
Requirement already satisfied: widgetsnbextension~3.5.0 in /usr/local/lib/python3.6/
Requirement already satisfied: webencodings in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: six>=1.9.0 in /usr/local/lib/python3.6/dist-packages (-
Requirement already satisfied: packaging in /usr/local/lib/python3.6/dist-packages (f
Collecting ray>=1.0.0; extra == "ray"

Downloading <https://files.pythonhosted.org/packages/12/87/44476ad712acc1f7957cbf88d>
|████████████████████| 23.1MB 2.0MB/s

Collecting pyarrow==1.0; extra == "ray"

Downloading <https://files.pythonhosted.org/packages/a1/0a/a89de6d747c4698af128a4639>
|████████████████████| 17.2MB 125kB/s

Collecting locket

Downloading <https://files.pythonhosted.org/packages/d0/22/3c0f97614e0be8386542facb3>
 Requirement already satisfied: simplegeneric>0.8 in /usr/local/lib/python3.6/dist-packages (from locket==0.1.1)
 Requirement already satisfied: decorator in /usr/local/lib/python3.6/dist-packages (from locket==0.1.1)
 Requirement already satisfied: prompt-toolkit<2.0.0,>=1.0.4 in /usr/local/lib/python3.6/dist-packages (from locket==0.1.1)
 Requirement already satisfied: pygments in /usr/local/lib/python3.6/dist-packages (from locket==0.1.1)
 Requirement already satisfied: setuptools>=18.5 in /usr/local/lib/python3.6/dist-packages (from locket==0.1.1)
 Requirement already satisfied: pexpect; sys_platform != "win32" in /usr/local/lib/python3.6/dist-packages (from locket==0.1.1)
 Requirement already satisfied: pickleshare in /usr/local/lib/python3.6/dist-packages (from locket==0.1.1)
 Requirement already satisfied: jsonschema!=2.5.0,>=2.4 in /usr/local/lib/python3.6/dist-packages (from locket==0.1.1)
 Requirement already satisfied: jupyter-core in /usr/local/lib/python3.6/dist-packages (from locket==0.1.1)
 Requirement already satisfied: ipython-genutils in /usr/local/lib/python3.6/dist-packages (from locket==0.1.1)
 Requirement already satisfied: tornado>=4.0 in /usr/local/lib/python3.6/dist-packages (from locket==0.1.1)
 Requirement already satisfied: jupyter-client in /usr/local/lib/python3.6/dist-packages (from locket==0.1.1)
 Requirement already satisfied: notebook>=4.4.1 in /usr/local/lib/python3.6/dist-packages (from locket==0.1.1)
 Requirement already satisfied: pyparsing>=2.0.2 in /usr/local/lib/python3.6/dist-packages (from locket==0.1.1)
 Collecting aiohttp-cors
 Downloading <https://files.pythonhosted.org/packages/13/e7/e436a0c0eb5127d8b491a9b83>
 Requirement already satisfied: requests in /usr/local/lib/python3.6/dist-packages (from aiohttp-cors==0.7.0)
 Collecting redis<3.5.0,>=3.3.2

▼ Import required libraries

```
import time
import swifter
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from datetime import datetime
import calendar

from textblob import TextBlob
import text2emotion as te

import gensim
from gensim import corpora
import pyLDAvis
import pyLDAvis.gensim

from sklearn.svm import LinearSVC
from sklearn.naive_bayes import MultinomialNB
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression, SGDClassifier
from sklearn.feature_extraction.text import TfidfVectorizer, CountVectorizer
from sklearn.metrics import accuracy_score, f1_score, precision_score, recall_score, classification_report

import nltk
from nltk import FreqDist
nltk.downloader.download('vader_lexicon')
from nltk.sentiment import SentimentAnalyzer
```

```
from nltk.sentiment import SentimentAnalyzer
```

```
import spacy
from spacy.lang.en.stop_words import STOP_WORDS
from spacy.lang.en import English
```

```
import re
import gzip
import string
import unicode
from bs4 import BeautifulSoup
import plotly.graph_objects as go
from wordcloud import WordCloud, STOPWORDS
```

```
import sys
import heapq
```

```
sns.set()
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Unzipping corpora/stopwords.zip.
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data]   Unzipping tokenizers/punkt.zip.
[nltk_data] Downloading package wordnet to /root/nltk_data...
[nltk_data]   Unzipping corpora/wordnet.zip.
[nltk_data] Downloading package vader_lexicon to /root/nltk_data...
/usr/local/lib/python3.6/dist-packages/nltk/twitter/__init__.py:20: UserWarning: The tw
warnings.warn("The twython library has not been installed. ")
```

```
"""
We will ignore FutureWarning and DeprecationWarning
"""

import warnings
warnings.filterwarnings("ignore", category=FutureWarning)
warnings.filterwarnings("ignore", category=DeprecationWarning)

"""
We will ignore warnings
"""

warnings.filterwarnings("ignore")

if not sys.warnoptions:
    warnings.simplefilter("ignore")
program_start_time=time.time()
```

▼ Download dataset

```
!wget https://s3.amazonaws.com/amazon-reviews-pds/tsv/amazon_reviews_us_Mobile_Electronics_v1
```

```
--2020-11-28 23:22:09-- https://s3.amazonaws.com/amazon-reviews-pds/tsv/amazon\_reviews
Resolving s3.amazonaws.com (s3.amazonaws.com)... 52.216.141.134
Connecting to s3.amazonaws.com (s3.amazonaws.com)|52.216.141.134|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 22870508 (22M) [application/x-gzip]
Saving to: 'amazon_reviews_us_Mobile_Electronics_v1_00.tsv.gz'
```

```
amazon_reviews_us_M 100%[=====>] 21.81M 23.8MB/s in 0.9s
```

```
2020-11-28 23:22:10 (23.8 MB/s) - 'amazon_reviews_us_Mobile_Electronics_v1_00.tsv.gz' saved
```



▼ Read the file

```
%time
with gzip.open('amazon_reviews_us_Mobile_Electronics_v1_00.tsv.gz') as f:
    df = pd.read_csv(f, sep='\t', error_bad_lines=False)

df.head()

b'Skipping line 35246: expected 15 fields, saw 22\n'
b'Skipping line 87073: expected 15 fields, saw 22\n'
CPU times: user 1.11 s, sys: 112 ms, total: 1.23 s
Wall time: 1.24 s
```

▼ Data information

```
df.shape

(104852, 15)
```

The dataset originally contains 104852 rows and 15 columns

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 104852 entries, 0 to 104851
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   marketplace            104852 non-null object
1   customer_id            104852 non-null int64
2   review_id              104852 non-null object
3   product_id             104852 non-null object
4   product_parent         104852 non-null int64
5   product_title          104852 non-null object
6   product_category       104852 non-null object
```

```

7  star_rating      104850 non-null float64
8  helpful_votes    104850 non-null float64
9  total_votes      104850 non-null float64
10 vine            104850 non-null object
11 verified_purchase 104850 non-null object
12 review_headline   104848 non-null object
13 review_body       104849 non-null object
14 review_date       104850 non-null object
dtypes: float64(3), int64(2), object(10)
memory usage: 12.0+ MB

```

```
df.columns
```

```

Index(['marketplace', 'customer_id', 'review_id', 'product_id',
      'product_parent', 'product_title', 'product_category', 'star_rating',
      'helpful_votes', 'total_votes', 'vine', 'verified_purchase',
      'review_headline', 'review_body', 'review_date'],
      dtype='object')

```

Columns Description:

marketplace: Marketplace of the product

customer_id: ID of the reviewer

review_id: ID of the review

product_id: ID of the product

product _parent: ID of the product parent

product_title: Product name

product_category: Category of the product

star_rating: Rating of the product from 1 to 5, 1 being the lowest

helpful_votes: Helpful votes of the review

total_votes: Total votes of the review

vine: Indicator for vine review

verified_purchase: Indicates if the purchase is verified or not

review_headline: Customer review title

```
review_body: Customer review summary
```

```
review_date: Customer review date
```

```
df.isnull().sum()
```

```
marketplace      0
customer_id      0
review_id        0
product_id       0
product_parent   0
product_title    0
product_category 0
star_rating      2
helpful_votes    2
total_votes      2
vine             2
verified_purchase 2
review_headline  4
review_body      3
review_date      2
dtype: int64
```

We have few reviews that do not contain certain information such as star_rating, review_headline etc.

```
# Making a copy of the dataframe
data = df.copy()
```

```
data.shape
```

```
(104852, 15)
```

▼ Missing data

```
# Dropping rows with missing information
data.dropna(axis = 0, how = 'any', inplace = True)
data = data.reset_index(drop=True)
```

```
data.isnull().sum()
```

```
marketplace      0
customer_id      0
review_id        0
product_id       0
product_parent   0
product_title    0
```

```

product_category    0
star_rating         0
helpful_votes       0
total_votes         0
vine                0
verified_purchase   0
review_headline     0
review_body         0
review_date         0
dtype: int64

```

```
data.shape
```

```
(104847, 15)
```

The dataset now contains 104847 rows and 15 columns

```

# Unique Reviews
len(pd.unique(data['review_id']))

104847

```

All the reviews are unique

▼ Data Statistics

```
data.describe()
```

| | customer_id | product_parent | star_rating | helpful_votes | total_votes |
|--------------|--------------|----------------|---------------|---------------|---------------|
| count | 1.048470e+05 | 1.048470e+05 | 104847.000000 | 104847.000000 | 104847.000000 |
| mean | 2.793730e+07 | 5.015301e+08 | 3.763160 | 1.244032 | 1.615440 |
| std | 1.508714e+07 | 2.871676e+08 | 1.523537 | 7.070485 | 7.910005 |
| min | 1.007100e+04 | 5.352400e+04 | 1.000000 | 0.000000 | 0.000000 |
| 25% | 1.471380e+07 | 2.593731e+08 | 3.000000 | 0.000000 | 0.000000 |
| 50% | 2.650319e+07 | 4.939016e+08 | 4.000000 | 0.000000 | 0.000000 |
| 75% | 4.223490e+07 | 7.440083e+08 | 5.000000 | 1.000000 | 1.000000 |
| max | 5.309657e+07 | 9.999508e+08 | 5.000000 | 769.000000 | 791.000000 |

```

# Builds histogram and set the number of bins and fig size (width, height)
data.hist(bins=50, figsize=(20,15))
plt.show()

```



▼ Deriving additional information

```
# Converting review_date to datetime object to extract month, day & year
%%time
data['review_date'] = pd.to_datetime(data['review_date'], format='%Y-%m-%d')
```



```
CPU times: user 25.2 ms, sys: 2.02 ms, total: 27.3 ms
Wall time: 29.7 ms
```

```
# Extracting month, day and year
```

```
%%time
```

```
data['day'] = data['review_date'].apply(lambda r:r.day)
data['month'] = data['review_date'].apply(lambda r:r.month)
data['year'] = data['review_date'].apply(lambda r:r.year)
```

```
CPU times: user 1.46 s, sys: 25.9 ms, total: 1.49 s
Wall time: 1.49 s
```

```
data.head(2)
```

| | marketplace | customer_id | review_id | product_id | product_parent | product_title |
|---|-------------|-------------|----------------|------------|----------------|---------------------------------|
| 0 | US | 20422322 | R8MEA6IGAHO0B | B00MC4CED8 | 217304173 | Black DR600 |
| 1 | US | 40835037 | R31LOQ8JGLPRLK | B00OQMFG1Q | 137313254 | GENSSI GPS Two Smart PI Car Ala |

▼ Review Trend over time

```
%%time
```

```
f, axes = plt.subplots(2,2, figsize=(12,8))
```

```
# Yearly Reviews
```

```
yearly = data.groupby(['year'])['review_id'].count().reset_index()
yearly = yearly.rename(columns={'review_id':'Number of reviews'})
year_trend = sns.lineplot(x='year',y='Number of reviews',data=yearly, ax = axes[0,0])
year_trend.set_title('Number of reviews over years')
```

```
# Monthly Reviews
```

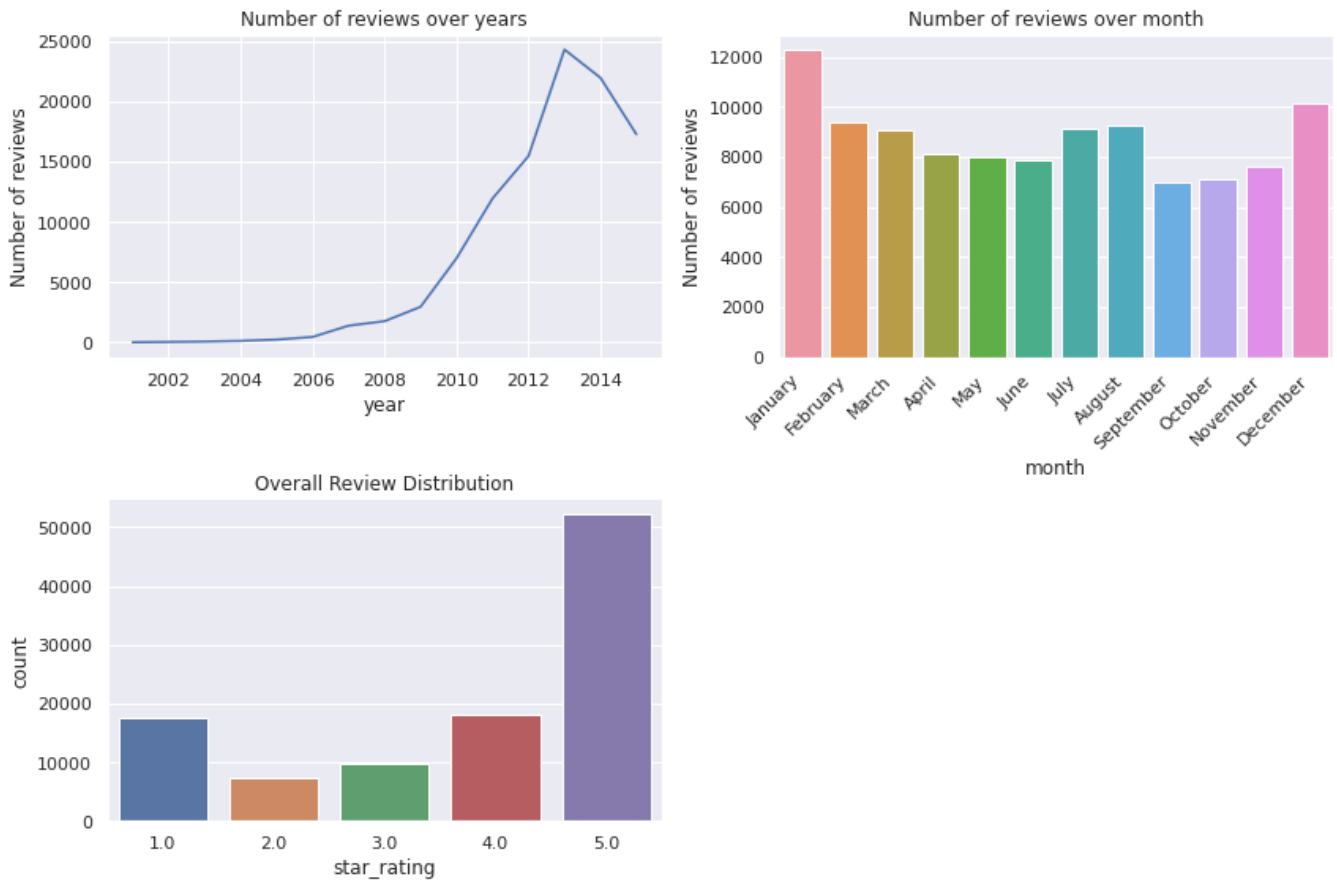
```
monthly = data.groupby(['month'])['review_id'].count().reset_index()
monthly['month'] = monthly['month'].apply(lambda x : calendar.month_name[x])
monthly = monthly.rename(columns={'review_id':'Number of reviews'})
month_trend = sns.barplot(x='month',y='Number of reviews',data=monthly, ax = axes[0,1])
month_trend.set_title('Number of reviews over month')
month_trend.set_xticklabels(month_trend.get_xticklabels(), rotation = 45, horizontalalignment
```

```
# Getting overall ratings for products
```

```
sns.countplot(x = 'star_rating', data = data, ax = axes[1,0] ).set_title('Overall Review Dist
```

```
f.delaxes(axes[1][1])
f.tight_layout()
```

CPU times: user 316 ms, sys: 91 ms, total: 407 ms
Wall time: 310 ms



Rating Trend over the years

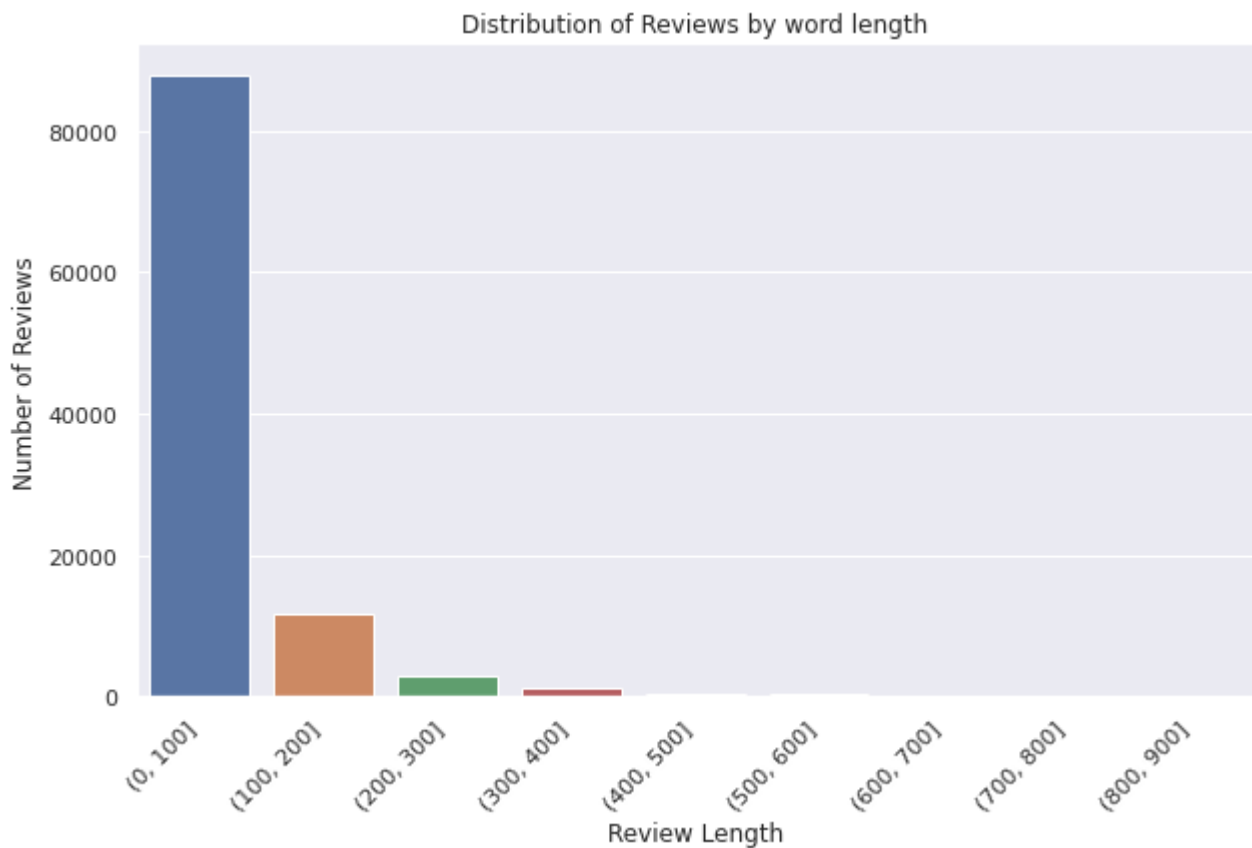
- There is an increasing trend for number of ratings given by the users to products on Amazon which indicates that a greater number of users started using the Amazon e-commerce site for online shopping and a greater number of users started giving feedback on the products purchased from 2006 to 2014. There is a significant increase in number of ratings given by users from 2012 to 2013.
- Notice the peak on 2013. Two major events support this. Amazon began to offer Sunday delivery option for purchases. See [news article](#) here. That surely resulted in lots of new members and increased ratings & reviews.

Distribution of overall ratings

- Many users have given a rating of 5 to products followed by 4 and 1 whereas very few users have given a low rating of 2 and 3.

▼ Distribution of reviews by word length

```
%%time
plt.figure(figsize=(10,6))
electronics_reviews = data[['review_id','customer_id','review_body','review_headline','star_r
electronics_reviews['review_length'] = electronics_reviews['review_body'].apply(lambda x: len
reviews_word_length = electronics_reviews.groupby(pd.cut(electronics_reviews.review_length, n
reviews_word_length = reviews_word_length.rename(columns={'review_length':'count'})
reviews_word_length = reviews_word_length.reset_index()
#print(reviews_word_length)
reviewLengthChart = sns.barplot(x='review_length',y='count',data=reviews_word_length)
reviewLengthChart.set_title('Distribution of Reviews by word length')
reviewLengthChart.set_xticklabels(reviewLengthChart.get_xticklabels(), rotation = 45, horizon
plt.xlabel("Review Length")
plt.ylabel("Number of Reviews")
plt.show()
```



CPU times: user 619 ms, sys: 4.61 ms, total: 624 ms
 Wall time: 631 ms

▼ Reviews per product

```
%%time
plt.figure(figsize=(10,6))
counts = data["product_title"].value_counts().to_frame()
counts.loc[counts['product_title'] > 250].plot(kind='bar')
plt.xlabel("Products")
plt.ylabel("Number of reviews")
plt.title("Number of reviews per product")
plt.show()
```



```
reviews['review_polarity'] = reviews['review'].swifter.apply(polarity)
```

Pandas Apply: 100%

104847/104847 [01:36<00:00, 1087.60it/s]


CPU times: user 1min 25s, sys: 680 ms, total: 1min 26s

Wall time: 1min 26s

```

# Plotting Histogram of Polarity Score
num_bins = 50
plt.figure(figsize=(10,6))
n, bins, patches = plt.hist(reviews.review_polarity, num_bins, facecolor='blue', alpha=0.5)
plt.xlabel('Polarity', fontsize=13)
plt.ylabel('Number of Reviews', fontsize=13)
plt.title('Histogram of Polarity Score', fontsize=13)
plt.show()

```



Although maximum reviews have positive emotions, there is peak for neutral reviews and be seen at 0.

```


# Calculating review subjectivity
%%time
sub = lambda x: TextBlob(x).sentiment.subjectivity
reviews['review_subjectivity'] = reviews['review'].swifter.apply(sub)

```

Pandas Apply: 100%


104847/104847 [02:12<00:00, 794.07it/s]

```
# Plotting distribution of subjectivity
plt.figure(figsize=(10,6))
sns.distplot(reviews.review_subjectivity, hist=True, kde=True, bins=int(30), color = 'darkblue')
plt.xlim([-0.001,1.001])
plt.xlabel('Subjectivity', fontsize=13)
plt.ylabel('Frequency', fontsize=13)
plt.title('Distribution of Subjectivity Score', fontsize=15)
plt.show()
```



There is quite a normal distribution for subjectivity. However, there are many reviews which are fully subjective and fully objective.

```
# Plotting polarity of star ratings
plt.figure(figsize=(10,6))
sns.boxenplot(x='star_rating', y='review_polarity', data=reviews)
plt.xlabel("Star Rating", fontsize=13)
plt.ylabel("Review Polarity", fontsize=13)
plt.title("Review Polarity of Star ratings", fontsize=13)
plt.show()
```



In the above plot, we can see that polarity increases with star rating. There are very few reviews with 5 rating and negative polarity and 1 rating with positive polarity.

```
#Some Positive reviews that has bad rating
```

```
reviews.loc[(reviews.review_polarity == -1) & (reviews.star_rating == 5)].review.head().tolist
```

```
['The mount that comes with the g1wh is horrible. Worked as described. The mount that co  
"THUMPER This sub pounds the hell out of my girls car, and for the price you can't go v  
"Crazy!!! Saw the video and bought this. You know this isn't Bose audio, but it works ;  
'terrible its does not record your voice for memos and you have to do it on the compute
```



All the above reviews have the negative words which might have made the polarity negative. These words are: horrible, hell, wrong, disappoint

```
#Some Negative reviews that has good rating
```

```
reviews.loc[(reviews.review_polarity == 1) & (reviews.star_rating == 1)].review.head().tolist
```

```
["One Star Wasn't very happy with it it did not have the power that I wanted",  
'One Star I am very happy with the product that I bought',  
"EARPOD doesn't work THE EARPOD DIDN'T WORK ( had to buy one by Best Buy)<br />DELIVERED  
'One Star The product had a shortage in the cord. I ordered another which worked perfect  
'One Star one works perfectly and one is a defective']
```



All the above reviews have the most positive words which might have made the polarity positive. These words are: Very happy, Best buy, perfectly

▼ Review lenath

```
# Calculating length of each review
```

```
length_of_review=[]
```

```
for i, word in enumerate(reviews.review.tolist()):
```

```
    word_length = len(word)
```

```
    length_of_review.append(word_length)
```

```
reviews['review_length'] =length_of_review
```

```
display(reviews)
```

| | review_headline | review_body | star_rating | review | review_polarity | review_s |
|-----|---|---|-------------|---|-----------------|----------|
| 0 | Very Happy! | As advertised. Everything works perfectly, I'm... | 5.0 | Very Happy! As advertised. Everything works pe... | 0.666667 | |
| 1 | five star | it's great | 5.0 | five star it's great | 0.800000 | |
| 2 | great cables | These work great and fit my life proof case fo... | 5.0 | great cables These work great and fit my life ... | 0.666667 | |
| 3 | Work very well but couldn't get used to not he... | Work very well but couldn't get used to not he... | 4.0 | Work very well but couldn't get used to not he... | 0.200000 | |
| 4 | Cameras has battery issues | Be careful with these products, I have bought ... | 2.0 | Cameras has battery issues Be careful with the... | 0.139225 | |
| ... | ... | ... | ... | ... | ... | |
| | The Cat Barf is | I've been looking for a | | The Cat Barf is Gone! I've | | |

```
# Plotting Star rating vs review length
```

```
L1= reviews[reviews['star_rating']==1]['review_length'].mean()
```

```
L2= reviews[reviews['star_rating']==2]['review_length'].mean()
```

```
L3= reviews[reviews['star_rating']==3]['review_length'].mean()
```

```
L4= reviews[reviews['star_rating']==4]['review_length'].mean()
```

```
L5= reviews[reviews['star_rating']==5]['review_length'].mean()
```


```
plt.figure(figsize=(10,6))
```

```
review_length_comparison=[L1,L2,L3,L4,L5]
```

```

review_length_comparison = [11, 22, 23, 21, 25]
objects = ('1 star', '2 star', '3 star', '4 star', '5 star')
y_pos = np.arange(len(objects))
plt.bar(y_pos, review_length_comparison, align='center', alpha=0.5)
plt.xticks(y_pos, objects)
plt.ylabel('Length of Review', fontsize=13)
plt.xlabel('Star Rating', fontsize=13)
plt.title('Review Length vs Rating', fontsize=13)
plt.show()

```



People tend to give very long reviews when they are explaining details about the product. For example, if someone doesn't like a particular feature of some product they explain it in detail. On the other hand, best reviews mostly are given in few words.

▼ Word Cloud

```


# Creating word cloud
def review_word_cloud(reviews):
    words = " ".join(reviews)

    wordcloud = WordCloud(stopwords=STOPWORDS,
                           background_color='black',
                           width=3000,
                           height=2500
                           ).generate(words)

```

```
plt.figure(1, figsize=(12, 12))
```

```
review word cloud(reviews["review"])
```



▼ Text2Emotion

Emotion Detection from text is useful in many ways. It helps to understand the customers, analyze feedback and reviews.

<https://colab.research.google.com/drive/1ws06lOR4-SUb3SpMejmPTn7JmajNGbs0#printMode=true>

```
t2e = t2e.review.map(lambda review, star_rating:
```

```
t2e = t2e.sample(frac = 0.1)
```

```
# Deriving emotion
start = time.time()
t2e["emotion_factor"] = t2e.review.map(te.get_emotion)
print((time.time()-start)/60, 'mins')
```

40.04022561709086 mins

```
%%time
t2e = pd.concat([t2e, t2e['emotion_factor'].apply(pd.Series)], axis = 1)
```

CPU times: user 3.44 s, sys: 121 ms, total: 3.57 s
Wall time: 3.49 s

t2e

| | review | star_rating | emotion_factor | Happy | Angry | Surprise | Sad | Fear |
|-------|--|-------------|---|-------|-------|----------|------|------|
| 70606 | Kick butt Nice amp for low watts. Running kenw... | 5.0 | {'Happy': 0.5, 'Angry': 0.0, 'Surprise': 0.25,... | 0.50 | 0.00 | 0.25 | 0.25 | 0.00 |
| 31226 | Really poor quality I got this to use at work ... | 1.0 | {'Happy': 0.13, 'Angry': 0.0, 'Surprise': 0.33... | 0.13 | 0.00 | 0.33 | 0.07 | 0.47 |
| 41292 | Great Price - All Necessary Pieces - Buy it Th... | 5.0 | {'Happy': 0.25, 'Angry': 0.0, 'Surprise': 0.38... | 0.25 | 0.00 | 0.38 | 0.25 | 0.12 |
| 95544 | Worthwhile Alternative For anyone who finds an... | 5.0 | {'Happy': 0.29, 'Angry': 0.0, 'Surprise': 0.0,... | 0.29 | 0.00 | 0.00 | 0.29 | 0.41 |
| 22203 | Very bad quality Broke in a month. After disas... | 2.0 | {'Happy': 0.0, 'Angry': 0.0, 'Surprise': 0.33,... | 0.00 | 0.00 | 0.33 | 0.67 | 0.00 |

```
# Deriving tone of the reviews
def get_tone(dct):
    sorted_d = sorted(dct.items(), key=lambda kv: kv[1], reverse=True)
    if sorted_d[1][1] != 0.0 and (sorted_d[0][1] != sorted_d[1][1]) :
        return 'More ' + sorted_d[0][0]+' Than '+sorted_d[1][0]
    elif sorted_d[0][1] == sorted_d[1][1]:
        return 'Both ' + sorted_d[0][0] + ' and ' + sorted_d[1][0]
    else:
        return sorted_d[0][0]
```

```
# Deriving tone of each review
t2e["Tone"] = t2e.emotion_factor.apply(get_tone)
t2e.head()
```

| | review | star_rating | emotion_factor | Happy | Angry | Surprise | Sad | Fear | Tone |
|--------------|---|-------------|--|-------|-------|----------|------|------|----------------------------|
| 70606 | Kick butt Nice amp for low watts. Running kenw... | 5.0 | {'Happy': 0.5, 'Angry': 0.0, 'Surprise': 0.25,... | 0.50 | 0.0 | 0.25 | 0.25 | 0.00 | Mo Hap Thi Surpri |
| 31226 | Really poor quality I got this to use at work ... Great | 1.0 | {'Happy': 0.13, 'Angry': 0.0, 'Surprise': 0.33... | 0.13 | 0.0 | 0.33 | 0.07 | 0.47 | Mo Fe Thi Surpri |

```
# Plotting emotions based on each rating
ll = t2e.groupby("star_rating", as_index=True)[['Happy', 'Angry', 'Sad', 'Surprise', 'Fear']].
ll.reset_index(inplace=True)

fig, axes = plt.subplots(2, 3, figsize=(15, 8))
plt.suptitle('Emotions based on each rating')

for i, (idx, row) in enumerate(ll.set_index('star_rating').iterrows()):
    ax = axes[i // 3, i % 3]
    row = row[row.gt(row.sum() * .01)]
    ax.pie(row, labels=row.index, startangle=30)
    ax.set_title(idx)
fig.delaxes(axes[1][2])
plt.show()
```

Emotions based on each rating



▼ Named Entity Recognition

NER helps to get a glance, understand the subject or theme of the reviews

```
# Getting top 20 most helpful votes
helpful = data["helpful_votes"].tolist()
most_helpful = heapq.nlargest(20, helpful)
```

```
# Joined the most helpful reviews
df_ner = data.loc[data.index.intersection(most_helpful)]
helpfultext = " ".join(df_ner['review_body'])
helpfultext
```

'Works great! Got it in red and black and already in love with it. Hands down one of the best Bluetooth speakers in this price range. If you're looking for cheap Bluetooth speakers SoundBot SB571 is the way to go. This is the worst radio I ever bought touch screen stop working very fast and the seller is the worst people to deal with don't buy from Waiter a month maybe longer, when it finally arrived it was not a Acten but some other model.
rear view camera works fine but navigation icon does not respond. Does not flow easily will return. It was loud at first and then after a couple weeks it kinda just gave up and wouldn't turn up even with the volume on my phone all the way up. And then the side of i

```
# Deriving entities
from collections import defaultdict
ner = spacy.load("en")
ner_helpful = ner(helpfultext)
ner_dict = defaultdict(list)
for entity in ner_helpful.ents:
    ner_dict[entity.label_].append(entity)
for NER, name in ner_dict.items():
    print(f"{NER}:\n{name}\n")

CARDINAL:
[one, two, two, about 10-15 feet, One, five]

ORG:
[SoundBot SB571, Sony, Sony, iPod, Considering, Toyota]
```

ORDINAL:
[first, first, first, first]

DATE:
[a couple weeks, a couple weeks, 2 days, two weeks]

GPE:
[kinda, kinda]

NORP:
[chinese]

MONEY:
[100 plus dollars]

LAW:
[the Camry Visor Repair]

```
# Getting the most reviewed product
product = data[data["product_id"] == 'B00J46X09U']
```

```
# Joining reviews
product_text = " ".join(product['review_body'])
product_text
```

'Very good quality. So far so good. Good product and good seller Great product! They charge my wife's phone Works like a charm! I've ordered a ton of these white and black, long and short. I keep buying them because they are made so well. I need more to buy for the office, my car, the house...etc... awesome stuff! It broke less than a month. Don't buy them. I've bought a total of 3 of these cables I've had the first for over a month and works great. Lasted longer than the ones I got from Apple. I needed one that worked with my lifeproof case and this does This cable works lightening-fast. As soon as I plug it into my iPhone, it recognizes the charger instantly and begins charging immediately. No

```
# Deriving entities
ner = spacy.load("en")
ner_helpful = ner(product_text)
ner_dict = defaultdict(list)
for entity in ner_helpful.ents:
    ner_dict[entity.label_].append(entity)
for NER, name in ner_dict.items():
    print(f"{NER}:\n{name}\n")
```

QUANTITY:
[a ton, 3ft, two 3 foot, 3 foot and, 6 foot, 3 foot, 6 foot, 2 charger, 6 ft cables, 3 ft]

DATE:
[less than a month, over a month, over a year, the day, almost three months, 2 days, dai]

CARDINAL:
[3, zero, two, one, 3, 6, 4, 1, one, 2, Only 1, 3, 6, 4, 2, 10, one, two, One, 6, 3, one]

ORDINAL:

[first, secondary, second, second, first, second, second, second, first, second, second,

ORG:

[Apple, iPhone, Apple, OEM, Apple, iPhone/ iPad Air, iPhone, Amazon, iPhone, USB, Apple,

MONEY:

[34;cheaper", 34, 34;skin", #6 iPhone, 34;accessory, 34;not, 34;certified", 34;

GPE:

[Lightning, Walmart, Otterbox, Apple", America, iPhone, iTunes, Lightning, China, Chi

NORP:

[Lifeproof, Lifeproof, Amazons, Working, Lightning, Lightning, Lightning]

PRODUCT:

[non-Apple, iPad, Amazon Basics, iPad, iPad, non-OEM, Chromo, non-Apple, iPads, USB, 6,

PERSON:

[Apple MFI, Excellent Cables, Nice Cables Great, jack, Charger, Lifeproof, Bummed, Light

LOC:

[iPhones, iPhones, iPhones, iPhones, the Amazon Basics, iPhones, SIX, the iPad Retina, i

FAC:

[5C, MacBook, Lightning, the Apple Store, 5Cs]

PERCENT:

[around 40%, 1%, 25% to 23%, 100%, 5 & 6 plus, 100%, 99%, 100%, 70%, 50%, 100%]

WORK_OF_ART:

[Apple MFI Certified, out=, Geniuses]

TIME:


[a few seconds, last night, 24 hours, midmorning, the afternoon, all night, last night,

LAW:

[the iPhone 6]

EVENT:

[iPhone, iPhone, iPhone]



```
# Using displacy to view entities
from spacy import displacy
displacy.render(ner_helpful, style="ent", jupyter=True)
```


Very good quality. So far so good. Good product and good seller Great product! They charge my wife's phone

Works like a charm! I've ordered a ton **QUANTITY** of these white and black, long and short. I keep buying them because they are made so well. I need more to buy for the office, my car, the house...etc... awesome st

It broke less than a month **DATE** . Don't buy them. I've bought a total of 3 **CARDINAL** of these cables

I've had the first **ORDINAL** for over a month **DATE** and works great. Lasted longer than the ones I got from Apple **ORG** . I needed one that worked with my lifeproof case and this does This cable works

lightening-fast. As soon as I plug it into my iPhone **ORG** , it recognizes the charger instantly and begins charging immediately. No need to wiggle the cable around the port to get it to fit properly or for the device to pick it up. I actually think it works better than the charger my iphone came with! I've had it for over a year

DATE now and the wire is still in excellent shape. No tears to the cable, like you often have with apple headphones or other chargers, plus the wire is still neat and clean hasn't really accumulated any dirt or smudges (I bought the white one so that says a lot). The cable also stays straight, i.e. it also doesn't tangle o

twist into a knot that you need to untangle. Whether you're looking for a primary iphone charger or a secondary **ORDINAL** charging cable for your car or for traveling, this product gets the job done!

Overall, a great item at a great price (no need to overpay at the Apple **ORG** store). It works very well ar

still lasts and looks pretty much the same as the day **DATE** I bought it. Well done with the manufacturing, definitely no knock-off experience here! Had this for almost three months **DATE** and it started to act up. There was no visual signs of any broken issues. The cord would charge the phone and then randomly stop charging. Worked just fine with my kids' iphones. Excellent product in a good price. You can't ask for anything better than that in retail. awesome product, not a thin flimsy cable like the OEM **ORG** part from Apple

ORG . This is very durable and has zero **CARDINAL** connectivity issues. I should have bought this a lon

time ago, and the two **CARDINAL** pack price is great. I have ordered these twice now (because my kids keep losing them !) and they hold up better than any others and charge the phone quickly Works great! Quali

is ok, but the cables tend to fray around the connection end fairly quickly. I have used several of these and lo

them. I have used a different #4 - 24-gauge #424 - MONY - brand before but you'll soon receive the error

▼ Text Preprocessing

```
nlp = spacy.load('en_core_web_sm')
spacy stopwords = spacy.lang.en.stop words.STOP WORDS
```

and charge my apple devices greatlv. Thank you! hard to find Lightning **GPE** cables that work for a long

```
def decontracted(phrase):
```

```
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)
    phrase = re.sub(r"n't", " not", phrase)
    phrase = re.sub(r"\ 're", " are", phrase)
    phrase = re.sub(r"\ 's", " is", phrase)
    phrase = re.sub(r"\ 'd", " would", phrase)
    phrase = re.sub(r"\ 'll", " will", phrase)
    phrase = re.sub(r"\ 't", " not", phrase)
    phrase = re.sub(r"\ 've", " have", phrase)
    phrase = re.sub(r"\ 'm", " am", phrase)
    phrase = re.sub(r" v ", " very", phrase)
    phrase = re.sub(r'\bthats\b', 'that is', phrase)
    phrase = re.sub(r'\bive\b', 'i have', phrase)
    phrase = re.sub(r'\bim\b', 'i am', phrase)
    phrase = re.sub(r'\bya\b', 'yeah', phrase)
    phrase = re.sub(r'\bcant\b', 'can not', phrase)
    phrase = re.sub(r'\bdont\b', 'do not', phrase)
    phrase = re.sub(r'\bwont\b', 'will not', phrase)
    phrase = re.sub(r'\bid\b', 'i would', phrase)
    phrase = re.sub(r'wtf', 'what the fuck', phrase)
    phrase = re.sub(r'\bwth\b', 'what the hell', phrase)
    phrase = re.sub(r'\br\b', 'are', phrase)
    phrase = re.sub(r'\bu\b', 'you', phrase)
    phrase = re.sub(r'\bk\b', 'OK', phrase)
    phrase = re.sub(r'\bsux\b', 'sucks', phrase)
    phrase = re.sub(r'\bno+\b', 'no', phrase)
    phrase = re.sub(r'\bcoo+\b', 'cool', phrase)
    phrase = re.sub(r'rt\b', '', phrase)
    phrase = phrase.strip()
    #print("decontracted:",phrase)
    return phrase
```

recommend these for anyone. Thanks. So far so good: fast shipping & good product. (had them about 2

```
# exclude words from spacy stopwords list
```

```
deselect_stop_words = ['no', 'not']
```

```
for w in deselect_stop_words:
```

```
    nlp.vocab[w].is_stop = False
```

```
# exclude words from spacy stopwords list
```

```
select_stop_words = ['#']
```

```
for w in select_stop_words:
```

```
    nlp.vocab[w].is_stop = True
```

combusted, fallen apart, or stopped working. Just take care of your things and they'll take care of you too.<br

```
def strip_html_tags(text):
```

```
    """remove html tags from text"""
```

```
    soup = BeautifulSoup(text, "html.parser")
```

```
    stripped_text = soup.get_text(separator=" ")
```

```
    #print("strip_html_tags:", stripped_text)
```

```

return stripped_text

    and are more durable than other cables I have used. It has even resisted some wear from my cat chewing on

def remove_accented_chars(text):
    """remove accented characters from text, e.g. café"""
    text = unidecode.unidecode(text)
    #print("remove_accented_chars:", text)
    return text

    ones Received two CARDINAL of these. One CARDINAL fall apart right when opened. These are chea

def remove_extra_characters(text):
    """remove extra characters from text, e.g. aaaaawwwweeeesssssoooooommmeeee"""
    text = re.sub("(.)\\1{2,}", "\\1", text)
    #print("remove_extra_characters:", text)
    return text

def keep_alphabet_numbers(text):
    """keep only words and numbers in the text"""
    text = re.sub('[^A-Za-z0-9]+', ' ', text)
    return text

def remove_urls(text):
    """remove url from the text"""
    # remove hyperlinks
    text = re.sub(r'\w+:\/\/{2}[\d\w-]+(\.[\d\w-]+)*(?:\/(?:\/[\^\/s\/]*))*', '', text)
    return text

    recommend NICE ORG issue free performance I purchased 3 CARDINAL packs of chargers over the

def text_preprocessing(text):
    """preprocess text with default option set to true for all steps"""
    text = strip_html_tags(text)
    text = remove_urls(text)
    text = remove_accented_chars(text)
    text = decontracted(text)
    text = remove_extra_characters(text)
    text = keep_alphabet_numbers(text)
    text = text.lower()
    tokens = nlp(text)
    review_text = [word for word in tokens if not word.is_stop]
    review_text = [word.lemma_ for word in review_text]
    return " ".join(review_text)

    """
    """

# Applying text-preprocessing to all the reviews
start = time.time()
reviews['preprocessed_review'] = reviews['review'].swifter.apply(lambda x: text_preprocessing

print((time.time()-start)/60, 'mins')

```

Pandas Apply: 100%

104847/104847 [26:27<00:00, 66.04it/s]

```
# Resetting the index
reviews = reviews.reset_index(drop=True)

# Double check to ensure there are no null values. Good practice to ensure no null values
reviews.head()
```

| | review_headline | review_body | star_rating | review | review_polarity | review_subject |
|---|--------------------|---|-------------|---|-----------------|----------------|
| 0 | Very Happy! | As advertised. Everything works perfectly, I'm... | 5.0 | Very Happy! As advertised. Everything works pe... | 0.666667 | 0.6 |
| 1 | five star | it's great | 5.0 | five star it's great | 0.800000 | 0.7 |
| 2 | great cables | These work great and fit my life proof case fo... | 5.0 | great cables These work great and fit my life ... | 0.666667 | 0.6 |
| | Work very well but | Work very well but | | Work very well but | | |

```
review_word_cloud(reviews["preprocessed_review"])
```


| | review_headline | review_body | star_rating | review | review_polarity | review_subject |
|---|---|---|-------------|---|-----------------|----------------|
| 0 | Very Happy! | As advertised. Everything works perfectly, I'm... | 5.0 | Very Happy! As advertised. Everything works pe... | 0.666667 | 0.6 |
| 1 | five star | it's great | 5.0 | five star it's great | 0.800000 | 0.7 |
| 2 | great cables | These work great and fit my life proof case fo... | 5.0 | great cables These work great and fit my life ... | 0.666667 | 0.6 |
| 3 | Work very well but couldn't get used to not he... | Work very well but couldn't get used to not he... | 4.0 | Work very well but couldn't get used to not he... | 0.200000 | 0.4 |

```
# Deriving average positive, negative and neutral sentiment of the dataset
```

```
avg_pos_sentiment = reviews['positive_sentiment'].mean()
avg_neu_sentiment = reviews['neutral_sentiment'].mean()
avg_neg_sentiment = reviews['negative_sentiment'].mean()
print("Average Positive Sentiment of dataset:", avg_pos_sentiment)
print("Average Neutral Sentiment of dataset:", avg_neu_sentiment)
print("Average Negative Sentiment of dataset:", avg_neg_sentiment)
```

```
Average Positive Sentiment of dataset: 0.3263287647715188
```

```
Average Neutral Sentiment of dataset: 0.5806806394079025
```

```
Average Negative Sentiment of dataset: 0.09294456684502155
```

```
from sklearn.feature_extraction.text import TfidfVectorizer
```

▼ Topic Modeling using LDA

Topic Modeling is a process to automatically identify topics present in a text object and to derive hidden patterns exhibited by a text corpus. Topic Modeling enables consumers to quickly extract the key topics covered by the reviews without having to go through all of them. It also helps the sellers/retailers get consumer feedback in the form of topics (extracted from the consumer reviews).

```
tfidf = TfidfVectorizer(stop_words='english')
tfidf.fit(docs)
word_embeddings = tfidf.get_feature_names_out()
```

```
# function to plot most frequent terms
```

```
def freq_words(x, terms = 30):
    all_words = ' '.join([text for text in x])
    all_words = all_words.split()

    fdist = FreqDist(all_words)
    words_df = pd.DataFrame({'word':list(fdist.keys()), 'count':list(fdist.values())})
```

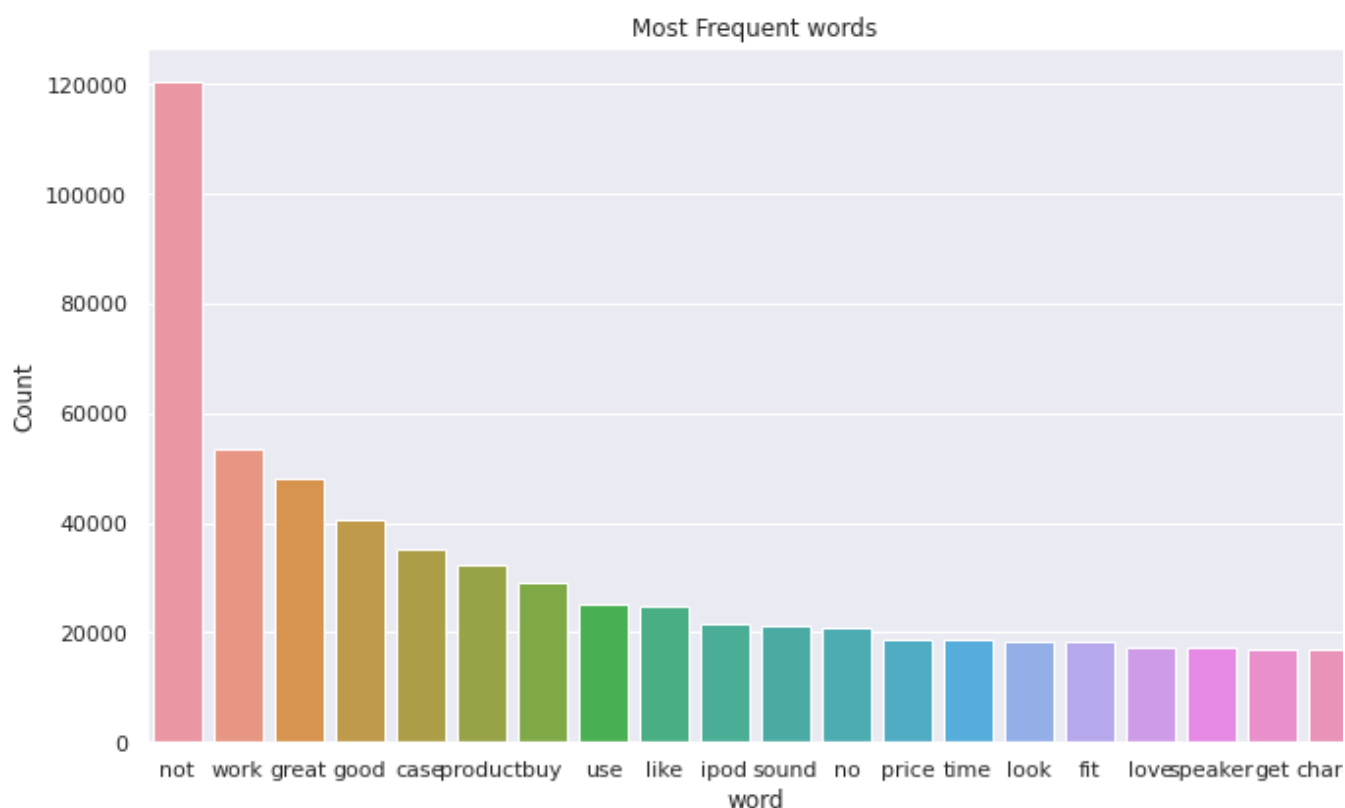
```
# selecting top 20 most frequent words
d = words_df.nlargest(columns="count", n = terms)
plt.figure(figsize=(10,6))
ax = sns.barplot(data=d, x= "word", y = "count")
ax.set(ylabel = 'Count')
plt.title("Most Frequent words")
plt.tight_layout()
plt.show()
```

```

# Creating the list of words from the reviews
review_list = [w for w in reviews['preprocessed_review']]

# Example:
manufactured cables that you can purchase from Apple ORG . Only been using for a month DATE

freq_words(review_list, 20)
```



```

Apple ORG cord may be worn and torn but it has never stopped charging. Your results may vary but I'm
tokenized_reviews = pd.Series(review_list).apply(lambda x: x.split())
print(tokenized_reviews[1])

['star', 'great']

or trying and that has not happened to this cord. It looks just like other lightning cables. I have iOS6 and a
dictionary = corpora.Dictionary(tokenized_reviews)
```

```

doc_term_matrix = [dictionary.doc2bow(rev) for rev in tokenized_reviews]

as well as the Apple ORG branded one CARDINAL Works great Have a iPhone 4s CARDINAL and

# Creating the object for LDA model using gensim library
lda = gensim.models.ldamodel.LdaModel
```

```
LDA = gensim.models.LdaModel.LdaModel
```

```
# Building LDA model
```

```
lda_model = LDA(corpus=doc_term_matrix, id2word=dictionary, num_topics=10, random_state=100,
                chunksize=1000, passes=25)
```

```
for a couple months DATE , not with it broken right stopped being recognized by my ...
```

```
# print topics
```

```
lda_model.print_topics()
```

```
[(0,
  '0.048*"scratch" + 0.040*"ok" + 0.031*"dock" + 0.028*"clean" + 0.026*"sub" + 0.022*"ch',
  (1,
    '0.041*"charge" + 0.031*"battery" + 0.031*"charger" + 0.029*"protector" + 0.024*"not"
  (2,
    '0.093*"not" + 0.024*"work" + 0.022*"buy" + 0.019*"product" + 0.013*"get" + 0.011*"tin',
  (3,
    '0.042*"player" + 0.033*"not" + 0.022*"ipod" + 0.021*"mp3" + 0.020*"use" + 0.019*"play',
  (4,
    '0.024*"screen" + 0.020*"video" + 0.015*"mount" + 0.014*"light" + 0.012*"software" + 0.011*',
  (5,
    '0.020*"gps" + 0.019*"unit" + 0.017*"book" + 0.015*"transmitter" + 0.015*"protection"
  (6,
    '0.081*"case" + 0.042*"ipod" + 0.034*"not" + 0.033*"fit" + 0.026*"cover" + 0.024*"scre',
  (7,
    '0.030*"radio" + 0.018*"unit" + 0.015*"car" + 0.014*"wire" + 0.010*"signal" + 0.010*"1',
  (8,
    '0.125*"great" + 0.089*"product" + 0.054*"good" + 0.051*"work" + 0.044*"price" + 0.022*',
  (9,
    '0.051*"sound" + 0.035*"speaker" + 0.028*"good" + 0.023*"great" + 0.022*"quality" + 0
```

since day one DATE . They do exactly what I want them to do. Excellent quality Works great to charge my

The Topic 1 has terms like 'charge', 'charger', 'battery' indicating that the topic is very much related to phone charging. Similarly, Topic 8 seems to be about the overall value of the product as it has terms like 'excellent', 'great', and 'recommend'

great and has been effectively charging the iPad PRODUCT . I like the length also because my daughter

To visualize our topics in a 2-dimensional space we used the pyLDAvis library. This visualization is interactive in nature and displays topics along with the most relevant words.

```
gained from the main screen. One ...
```

```
# Visualize the topics
```

```
pyLDAvis.enable_notebook()
```


```
vis = pyLDAvis.gensim.prepare(lda_model, doc_term_matrix, dictionary)
```

```
vis
```


Selected Topic:

Slid

Intertopic Distance Map (via multidimensional scaling)



▼ Classification Models

```
# Defining features and labels
review_data = reviews.copy()
y = review_data['star_rating'].values
y = y.astype('int')
X = review_data['preprocessed_review']
```

```
# Tfidf vectorizer
```

```

vectorizer = TfidfVectorizer(
    min_df=2,
    max_df=0.95,
    ngram_range = (1,4),
    stop_words = 'english',
)
# Extract features from reviews.
review_features = vectorizer.fit_transform(X)
    :) Got this cable a while ago but haven't used it heavily yet. I noticed it doesn't charge sometimes on one or tr

# Split the dataset to be 80% training and 20% testing
X_train, X_test, y_train, y_test = train_test_split(review_features,
                                                    y,
                                                    stratify=y,
                                                    random_state = 42,
                                                    test_size = 0.2
                                                    )

print(X_train.shape)

(83877, 520053)

```

```

# Logistic Regression model
lr = LogisticRegression(random_state=22).fit(X_train, y_train)
print( classification_report(y_test, lr.predict(X_test), digits=4))

```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 1 | 0.6684 | 0.8002 | 0.7284 | 3514 |
| 2 | 0.3813 | 0.0671 | 0.1142 | 1460 |
| 3 | 0.4792 | 0.2073 | 0.2894 | 1944 |
| 4 | 0.4304 | 0.2499 | 0.3162 | 3613 |
| 5 | 0.7125 | 0.9260 | 0.8054 | 10439 |
| accuracy | | | 0.6620 | 20970 |
| macro avg | 0.5344 | 0.4501 | 0.4507 | 20970 |
| weighted avg | 0.6118 | 0.6620 | 0.6122 | 20970 |

were as intended. Do not buy other cheap brand cables I have made that mistake in the past. Trust the cable

```

# MultinomialNB model
nb = MultinomialNB()
nb.fit(X_train, y_train)
print( classification_report(y_test, nb.predict(X_test), digits=4))

```

| | precision | recall | f1-score | support |
|-----------|-----------|--------|----------|---------|
| 1 | 0.8732 | 0.2783 | 0.4221 | 3514 |
| 2 | 0.0000 | 0.0000 | 0.0000 | 1460 |
| 3 | 0.7222 | 0.0067 | 0.0133 | 1944 |
| 4 | 0.6471 | 0.0030 | 0.0061 | 3613 |
| 5 | 0.5263 | 0.9989 | 0.6894 | 10439 |
| accuracy | | | 0.5451 | 20970 |
| macro avg | 0.5538 | 0.2574 | 0.2262 | 20970 |

| | | | | |
|--------------|--------|--------|--------|-------|
| weighted avg | 0.5867 | 0.5451 | 0.4162 | 20970 |
|--------------|--------|--------|--------|-------|

```

# SGDClassifier
sgd = SGDClassifier(random_state=22)
sgd.fit(X_train, y_train)
print( classification_report(y_test, sgd.predict(X_test), digits=4))

```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 1 | 0.6372 | 0.7866 | 0.7040 | 3514 |
| 2 | 0.2708 | 0.0089 | 0.0172 | 1460 |
| 3 | 0.5647 | 0.1055 | 0.1777 | 1944 |
| 4 | 0.4252 | 0.0551 | 0.0975 | 3613 |
| 5 | 0.6469 | 0.9762 | 0.7782 | 10439 |
| accuracy | | | 0.6377 | 20970 |
| macro avg | 0.5090 | 0.3864 | 0.3549 | 20970 |
| weighted avg | 0.5733 | 0.6377 | 0.5398 | 20970 |

phone while charging. Phone is basically unusable when plugged in. The page you are viewing bounces all

```

# RandomForestClassifier
forest = RandomForestClassifier(n_estimators=25, criterion="entropy", random_state=42)
forest.fit(X_train, y_train)
print( classification_report(y_test, forest.predict(X_test), digits=4))

```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 1 | 0.6490 | 0.6713 | 0.6600 | 3514 |
| 2 | 0.2812 | 0.0123 | 0.0236 | 1460 |
| 3 | 0.4367 | 0.0514 | 0.0920 | 1944 |
| 4 | 0.3326 | 0.0435 | 0.0769 | 3613 |
| 5 | 0.6101 | 0.9685 | 0.7486 | 10439 |
| accuracy | | | 0.6077 | 20970 |
| macro avg | 0.4619 | 0.3494 | 0.3202 | 20970 |
| weighted avg | 0.5299 | 0.6077 | 0.5067 | 20970 |

```

# LinearSVC
svc = LinearSVC(C = 20, class_weight= 'balanced')
svc.fit(X_train, y_train)
print( classification_report(y_test, svc.predict(X_test), digits=4))

```

| | precision | recall | f1-score | support |
|----------|-----------|--------|----------|---------|
| 1 | 0.6602 | 0.7271 | 0.6920 | 3514 |
| 2 | 0.2690 | 0.1555 | 0.1970 | 1460 |
| 3 | 0.3235 | 0.2320 | 0.2702 | 1944 |
| 4 | 0.3577 | 0.3006 | 0.3267 | 3613 |
| 5 | 0.7401 | 0.8384 | 0.7862 | 10439 |
| accuracy | | | 0.6233 | 20970 |

| | | | | |
|--------------|--------|--------|--------|-------|
| macro avg | 0.4701 | 0.4507 | 0.4544 | 20970 |
| weighted avg | 0.5894 | 0.6233 | 0.6024 | 20970 |

CARDINAL pack and it was a great price and they do the job well. Not like some less expensive cords, these

▼ Model Tuning

We performed parameter tuning using sklearn's GridSearchCV algorithm. GridSearchCV does an exhaustive search over specified parameter values.

```
from sklearn.model_selection import GridSearchCV

bought me 2 CARDINAL pack. One CARDINAL cable works very well in my car and hasn't failed me as

parameters = { 'alpha': [0.1, 0.25, 0.5, 0.75, 1.0, 1.25, 1.5, 1.75, 2] }
nb_clf = GridSearchCV(MultinomialNB(), parameters)
nb_clf.fit(X_train, y_train)

GridSearchCV(cv=None, error_score=nan,
             estimator=MultinomialNB(alpha=1.0, class_prior=None,
                                     fit_prior=True),
             iid='deprecated', n_jobs=None,
             param_grid={'alpha': [0.1, 0.25, 0.5, 0.75, 1.0, 1.25, 1.5, 1.75,
                                   2]},
             pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
             scoring=None, verbose=0)

on these cords and have broken others in just days DATE . Price is low in comparison. I've ordered 2

# Print out the parameters GridSearchCV decided on.
nb_clf.best_params_

{'alpha': 0.1}

a few weeks DATE we are happy with them. However, these cords would not charge a friend's iphone 5

parameters = [{ 'loss': ['hinge', 'log', 'perceptron'],
                 'alpha': 10.0**-np.arange(1,7),
                 'penalty': ['l1', 'l2', 'elasticnet']}]
sgd_clf = GridSearchCV(SGDClassifier(random_state=22), parameters)
sgd_clf.fit(X_train, y_train)

GridSearchCV(cv=None, error_score=nan,
             estimator=SGDClassifier(alpha=0.0001, average=False,
                                     class_weight=None, early_stopping=False,
                                     epsilon=0.1, eta0=0.0, fit_intercept=True,
                                     l1_ratio=0.15, learning_rate='optimal',
                                     loss='hinge', max_iter=1000,
                                     n_iter_no_change=5, n_jobs=None,
                                     penalty='l2', power_t=0.5, random_state=22,
                                     shuffle=True, tol=0.001,
                                     validation_fraction=0.1, verbose=0,
                                     warm_start=False),
             iid='deprecated', n_jobs=None,
             param_grid=[{'alpha': array([1.e-01, 1.e-02, 1.e-03, 1.e-04, 1.e-05, 1.e-06,
```

```

        'loss': ['hinge', 'log', 'perceptron'],
        'penalty': ['l1', 'l2', 'elasticnet']]],
    pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
    scoring=None, verbose=0)

```

wires.....but it just stopped working. we'll see how the other one does, but it's not looking so promising. no

```
# Print out the parameters GridSearchCV decided on.
```

```
sgd_clf.best_params_
```

```
{'alpha': 1e-05, 'loss': 'log', 'penalty': 'l1'}
```

went through wash & dry **ORG** . Still works great! Can't beat that, include a 1 year **DATE** warranty!

```
classifiers = [lr, nb, sgd, forest, svc]
```

```
names_of_classifiers = ["Logistic Regression", "Multinomial Naive Bayes", "SGD", "Random Forest
```

cases for it none o/o plus as described, worked w/out error msg & charges quickly! very good quality and wa

```
# Create, fit and predict the star ratings using the Machine Learning classifiers
```

```
acc_score=[]
```

```
for index,clf in enumerate(classifiers):
```

```
    y_pred = clf.predict(X_test)
```

```
    acc= accuracy_score(y_test, y_pred)
```

```
    acc_score.append(acc*100)
```

```
if index == len(classifiers) - 1:
```

```
    cm=confusion_matrix(y_test,y_pred)
```

```
import os
```

CARDINAL stars. Works great so far. These are Apple **ORG** Certified! DO NOT BUY CORDS FOR YOU

```
# Plotting the accuracy comparison
```

```
x = names_of_classifiers
```

```
acc = acc_score
```

```
fig, ax = plt.subplots()
```

```
width = 0.75 # the width of the bars
```

```
ind = np.arange(len(acc)) # the x locations for the groups
```

```
ax.barh(ind, acc, width, color=(0.2, 0.4, 0.6, 0.6))
```

```
ax.set_yticks(ind+width/2)
```

```
ax.set_yticklabels(x, minor=False)
```

```
for i, v in enumerate(acc):
```

```
    temp = float(v)
```

```
    ax.text(0.5, i+0.1, str(round(v,2)))
```

```
plt.title('Accuracy Comparision')
```

```
plt.xlabel('Accuracy')
```

```
plt.ylabel('Classifiers')
```

```
#plt.show()
```

```
plt.savefig(os.path.join('test.png'), dpi=300, format='png', bbox_inches='tight')
```

Package came on time- everything arrived as outlined in listing. Good! These actually work!! As expected

```
# print classification reports
```

```

print("\nclassification reports: ")
print( "LogisticRegression: " )
print( classification_report(y_test, lr.predict(X_test), digits=4))
print( "MultinomialNB: " )
print( classification_report(y_test, nb_clf.predict(X_test), digits=4))
print( "SGDClassifier: " )
print( classification_report(y_test, sgd_clf.predict(X_test), digits=4))
print( "Random Forest: " )
print( classification_report(y_test, forest.predict(X_test), digits=4))
print( "SVC: " )
print( classification_report(y_test, svc.predict(X_test), digits=4))

```

classification reports:

LogisticRegression:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 1 | 0.6684 | 0.8002 | 0.7284 | 3514 |
| 2 | 0.3813 | 0.0671 | 0.1142 | 1460 |
| 3 | 0.4792 | 0.2073 | 0.2894 | 1944 |
| 4 | 0.4304 | 0.2499 | 0.3162 | 3613 |
| 5 | 0.7125 | 0.9260 | 0.8054 | 10439 |
| accuracy | | | 0.6620 | 20970 |
| macro avg | 0.5344 | 0.4501 | 0.4507 | 20970 |
| weighted avg | 0.6118 | 0.6620 | 0.6122 | 20970 |

MultinomialNB:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 1 | 0.6479 | 0.7610 | 0.6999 | 3514 |
| 2 | 0.3100 | 0.0212 | 0.0397 | 1460 |
| 3 | 0.4005 | 0.0880 | 0.1442 | 1944 |
| 4 | 0.3522 | 0.2380 | 0.2841 | 3613 |
| 5 | 0.6855 | 0.9111 | 0.7824 | 10439 |
| accuracy | | | 0.6317 | 20970 |
| macro avg | 0.4792 | 0.4039 | 0.3901 | 20970 |
| weighted avg | 0.5692 | 0.6317 | 0.5718 | 20970 |

SGDClassifier:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 1 | 0.6588 | 0.8011 | 0.7230 | 3514 |
| 2 | 0.3309 | 0.0315 | 0.0575 | 1460 |
| 3 | 0.5050 | 0.2068 | 0.2934 | 1944 |
| 4 | 0.4511 | 0.2209 | 0.2965 | 3613 |
| 5 | 0.7025 | 0.9417 | 0.8047 | 10439 |
| accuracy | | | 0.6624 | 20970 |
| macro avg | 0.5297 | 0.4404 | 0.4350 | 20970 |
| weighted avg | 0.6077 | 0.6624 | 0.6040 | 20970 |

Random Forest:

| | precision | recall | f1-score | support |
|--|-----------|--------|----------|---------|
|--|-----------|--------|----------|---------|

| | | | | |
|---|--------|--------|--------|-------|
| 1 | 0.6490 | 0.6713 | 0.6600 | 3514 |
| 2 | 0.2812 | 0.0123 | 0.0236 | 1460 |
| 3 | 0.4367 | 0.0514 | 0.0920 | 1944 |
| 4 | 0.3326 | 0.0435 | 0.0769 | 3613 |
| 5 | 0.6101 | 0.9685 | 0.7486 | 10439 |

| | | | | |
|--------------|--------|--------|--------|-------|
| accuracy | | | 0.6077 | 20970 |
| macro avg | 0.4619 | 0.3494 | 0.3202 | 20970 |
| weighted avg | 0.5299 | 0.6077 | 0.5067 | 20970 |

SVC:

| | precision | recall | f1-score | support |
|---|-----------|--------|----------|---------|
| 1 | 0.6602 | 0.7271 | 0.6920 | 3514 |
| 2 | 0.2690 | 0.1555 | 0.1970 | 1460 |

```
# Create, fit and predict the star ratings using the Machine Learning classifiers
```

```
acc_score=[]
```

```
refined_classifiers = [lr, nb_clf, sgd_clf, forest, svc]
```

```
for index,clf in enumerate(refined_classifiers):
```

```
    y_pred = clf.predict(X_test)
```

```
    acc= accuracy_score(y_test, y_pred)
```

```
    acc_score.append(acc*100)
```

```
if index == len(refined_classifiers) - 1:
```

```
    cm=confusion_matrix(y_test,y_pred)
```

```
x = ["Logistic Regression","Multinomial Naive Bayes","SGD", "Random Forest", "Linear SVC"]
```

```
acc = acc_score
```

```
fig, ax = plt.subplots()
```

```
width = 0.75 # the width of the bars
```

```
ind = np.arange(len(acc)) # the x locations for the groups
```

```
ax.barh(ind, acc, width, color=(0.2, 0.4, 0.6, 0.6))
```

```
ax.set_yticks(ind+width/2)
```

```
ax.set_yticklabels(x, minor=False)
```

```
for i, v in enumerate(acc):
```

```
    temp = float(v)
```

```
    ax.text(0.5, i+0.1,str(round(v,2)))
```


```
plt.title('Accuracy Comparsion after model tuning')
```

```
plt.xlabel('Accuracy')
```

```
plt.ylabel('Classifiers')
```

```
#plt.show()
```

```
plt.savefig(os.path.join('test.png'), dpi=300, format='png', bbox_inches='tight')
```



Plotting confusion matrix

```
def plot_confusion_matrix(cm, title='Confusion matrix'):
    plt.figure( figsize=(9,4))
    plt.imshow(cm, interpolation='nearest')
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(set(y)))
    plt.xticks(tick_marks, set(y), rotation=45)
    plt.yticks(tick_marks, set(y))
    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label')
```


rarely used has already stopped working. Not happy with this product. Works! Great and much cheaper

```
def create_and_print_confusion_matrix(y_test, predicted, title):
    cm = confusion_matrix(y_test, predicted)
    np.set_printoptions(precision=2)
    plt.figure()
    plot_confusion_matrix(cm, title)
    plt.show()
```

bought 2 **CARDINAL** of these cables for an iphone 5s **CARDINAL** , and iphone 5.
Works great, ve

```
create_and_print_confusion_matrix(y_test, lr.predict(X_test), "LogisticRegression")
```

<Figure size 432x288 with 0 Axes>

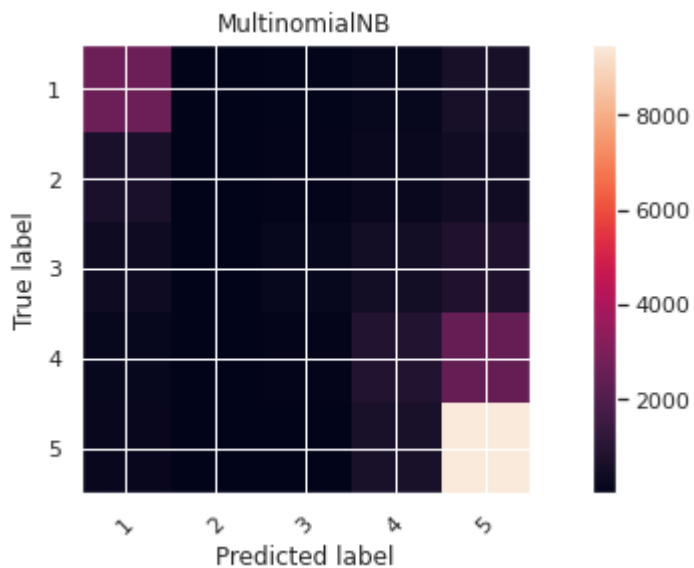


they don't last (they're the same junk cables you one can find on Amazon **ORG** for \$ 2-\$5 **MONEY**). I

Create MultinomialNB confusion matrix.

```
create_and_print_confusion_matrix(y_test, nb_clf.predict(X_test), "MultinomialNB")
```

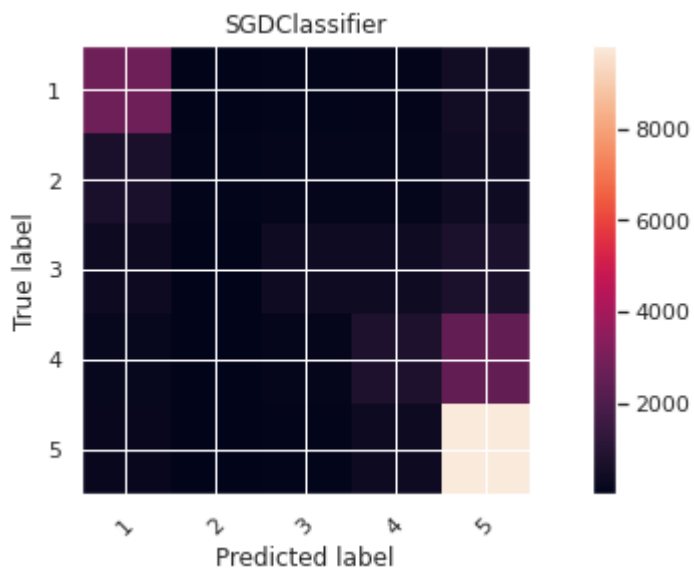

<Figure size 432x288 with 0 Axes>



Create SGDClassifier confusion matrix.


create_and_print_confusion_matrix(y_test, sgd_clf.predict(X_test), "SGDClassifier")

<Figure size 432x288 with 0 Axes>



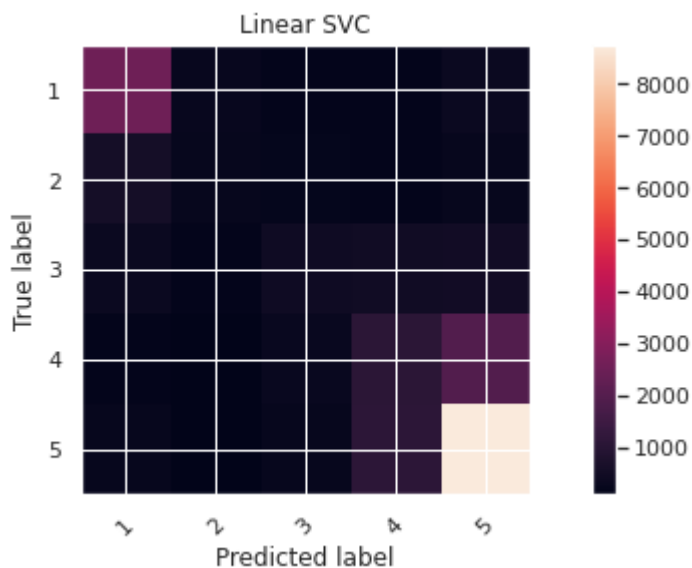
create_and_print_confusion_matrix(y_test, forest.predict(X_test), "Random Forest")

<Figure size 432x288 with 0 Axes>



```
create_and_print_confusion_matrix(y_test, svc.predict(X_test), "Linear SVC")
```

<Figure size 432x288 with 0 Axes>



REFERENCES:

Band, A. (2020, September 27). Text2emotion: Python package to detect emotions from textual data. Retrieved November 29, 2020, from <https://towardsdatascience.com/text2emotion-python-package-to-detect-emotions-from-textual-data-b2e7b7ce1153>

Confusion Matrix in Machine Learning. (2020, August 21). Retrieved November 29, 2020, from <https://www.geeksforgeeks.org/confusion-matrix-machine-learning/>

Joshi, P. (2020, June 25). SpaCy Tutorial: SpaCy For NLP: SpaCy NLP Tutorial. Retrieved November 29, 2020, from <https://www.analyticsvidhya.com/blog/2020/03/spacy-tutorial-learn-natural-language-processing/>

Kadam, S. (2020, April 18). Generating Word Cloud in Python. Retrieved November 29, 2020, from <https://www.geeksforgeeks.org/generating-word-cloud-python/>

Logistic Regression 3-class Classifier¶. (n.d.). Retrieved November 29, 2020, from https://scikit-learn.org/stable/auto_examples/linear_model/plot_iris_logistic.html?highlight=logistic+regression

Notebook on nbviewer. (n.d.). Retrieved November 29, 2020, from https://nbviewer.jupyter.org/github/bmabey/hacker_news_topic_modelling/blob/master/HN Topic

Model Talk.ipynb

Rai, A. (2019, January 23). Python: Sentiment Analysis using VADER. Retrieved November 29, 2020, from <https://www.geeksforgeeks.org/python-sentiment-analysis-using-vader/>

Simplified Text Processing¶. (n.d.). Retrieved November 29, 2020, from <https://textblob.readthedocs.io/en/dev/>

Sklearn.model_selection.GridSearchCV¶. (n.d.). Retrieved November 29, 2020, from https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.GridSearchCV.html

Only had this set a few weeks DATE , but seem sturdier than other brands I have tried. Also comes with a 12 month DATE warranty i was not aware of, which is good because it seems like the replacement ones I bought in the past wear out far sooner than that.
SIDE NOTE: LIKE THIS IF YOU WANT APPLE TO CHANGE THEIR CHARGERS. I Love apple, but I soon be switching if these new chargers continue to fail so quickly, but until then, these seem to be a good quality cable Good price, fit my case (which most non Apple ORG chargers doesnt) and no problems in a month DATE of using it daily I tried a lot of different "alternate" MONEY ; iPhone 5 CARDINAL cords and all failed - until these - I love them! They're great, solid product! Seems to work really well with my i Devices. No "Compatibility" MONEY ; messages :-) Well made and durable so far. We have 2 CARDINAL iPhone ORG 6's and and iPad. So far these cords work great on all the devices. Great price for a good product. Will buy again if/when we need additional cords. Well made cable. I give it 4 CARDINAL instead of 5 CARDINAL because after using it for about 3 weeks DATE my iPhone 6 CARDINAL is now telling me it is not a compatible device. I